

AlcoholConsumptionAnalysis2021

September 17, 2021

1 Exploratory Data Analysis: Alcohol Consumption Across the World

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Below, I imported two csv files to merge into one new dataframe: 1. `alcoholdata` dataframe containing **alcohol consumption across the world** 2. `countrycontinentmatch` dataframe **maps the country name to the continent**.

1.0.1 Data Overview

I wanted to see what types of data each dataframe comprises of using: - `df.head()` to view the first 5 rows of the dataframe - `df.dtypes` to view the types of data in each column

```
[2]: alcoholdata=pd.read_csv('https://raw.githubusercontent.com/fivethirtyeight/data/
↳master/alcohol-consumption/drinks.csv')
alcoholdata.head()
```

```
[2]:      country  beer_servings  spirit_servings  wine_servings  \
0  Afghanistan           0           0           0
1    Albania           89          132           54
2    Algeria           25           0           14
3   Andorra          245          138          312
4    Angola           217           57           45

      total_litres_of_pure_alcohol
0                0.0
1                4.9
2                0.7
3             12.4
4                5.9
```

```
[3]: alcoholdata.dtypes
```

```
[3]: country          object
     beer_servings    int64
     spirit_servings   int64
     wine_servings     int64
     total_litres_of_pure_alcohol float64
     dtype: object
```

```
[4]: countrycontinentmatch=pd.read_csv('https://pkgstore.datahub.io/JohnSnowLabs/
    ↪country-and-continent-codes-list/country-and-continent-codes-list-csv_csv/
    ↪data/b7876b7f496677669644f3d1069d3121/
    ↪country-and-continent-codes-list-csv_csv.csv')
     countrycontinentmatch.head()
```

```
[4]:  Continent_Name Continent_Code Country_Name \
0      Asia          AS      Afghanistan, Islamic Republic of
1      Europe         EU      Albania, Republic of
2      Antarctica      AN  Antarctica (the territory South of 60 deg S)
3      Africa          AF      Algeria, People's Democratic Republic of
4      Oceania         OC      American Samoa

     Two_Letter_Country_Code Three_Letter_Country_Code Country_Number
0                        AF          AFG          4.0
1                        AL          ALB          8.0
2                        AQ          ATA          10.0
3                        DZ          DZA          12.0
4                        AS          ASM          16.0
```

1.0.2 Data Manipulation: Extracting the country name

I only need the first word of the **Country_Name** column, thus I will use the `split()` method to return a list of all the words in the string, and extract the first word, followed by replacing the commas with a whitespace. - `str.split()`: This returns a list of lines - `str.split().str[x]`: This returns a xth value in the string - `str.replace(['old value'],'new value')`: This replaces values in a dataframe

```
[5]: countrycontinentmatch.dtypes
```

```
[5]: Continent_Name          object
     Continent_Code         object
     Country_Name           object
     Two_Letter_Country_Code object
     Three_Letter_Country_Code object
     Country_Number         float64
     dtype: object
```

```
[6]:
```

```
#splitting the 'Country_Name' column into lines of lists and extracting the
↳first word with str[0]
countrycontinentmatch['Name']=countrycontinentmatch['Country_Name'].str.split('
↳').str[0]
countrycontinentmatch['Name']
```

```
[6]: 0      Afghanistan,
      1      Albania,
      2      Antarctica
      3      Algeria,
      4      American
      ...
      257     Zambia,
      258     Disputed
      259     Iraq-Saudi
      260     United
      261     Spratly
      Name: Name, Length: 262, dtype: object
```

```
[7]: #removing the commas from the name column
countrycontinentmatch['Name']=countrycontinentmatch['Name'].str.replace(',','')
countrycontinentmatch['Name']
```

```
[7]: 0      Afghanistan
      1      Albania
      2      Antarctica
      3      Algeria
      4      American
      ...
      257     Zambia
      258     Disputed
      259     Iraq-Saudi
      260     United
      261     Spratly
      Name: Name, Length: 262, dtype: object
```

```
[8]: countrycontinentmatch.head()
```

```
[8]:  Continent_Name  Continent_Code  Country_Name \
0      Asia        AS      Afghanistan, Islamic Republic of
1      Europe      EU      Albania, Republic of
2      Antarctica  AN  Antarctica (the territory South of 60 deg S)
3      Africa      AF      Algeria, People's Democratic Republic of
4      Oceania     OC      American Samoa

      Two_Letter_Country_Code  Three_Letter_Country_Code  Country_Number \
0      AF                    AFG                    4.0
```

1	AL	ALB	8.0
2	AQ	ATA	10.0
3	DZ	DZA	12.0
4	AS	ASM	16.0

	Name
0	Afghanistan
1	Albania
2	Antarctica
3	Algeria
4	American

1.0.3 Data Merging: Merging Two Datasets with a Common Column Name

Here, I want to classify the countries in the **alcoholdata** dataframe to match the continents in the **countrycontinentmatch** dataframe.

I am only interested in the continent name and country for the **countrycontinentmatch** dataframe, let's extract that into its own dataframe **countires**. In this new dataframe, it only consists of 2 columns: Continent, country. Note that I renamed the name to country to match the column name in the alcoholdata dataframe, so that I can leverage the merge function to automatically join using *country* as a common key. - **pd.merge()**: Automatically joins dataframes using a common key. If not we can use *left_on* or *right_on* keywords to specify the different column names.

The output of this is a new dataframe **newalcoholdata** with an additional column for *Continent*

```
[9]: countries=countrycontinentmatch[['Continent_Name','Name']]
countries.rename(columns={'Continent_Name':'Continent','Name':
↳ 'country'},inplace=True)
countries
```

```
/opt/conda/envs/ADA522/lib/python3.8/site-packages/pandas/core/frame.py:4441:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
return super().rename(
```

```
[9]:
```

	Continent	country
0	Asia	Afghanistan
1	Europe	Albania
2	Antarctica	Antarctica
3	Africa	Algeria
4	Oceania	American
..
257	Africa	Zambia

```

258      Oceania      Disputed
259          Asia      Iraq-Saudi
260          Asia          United
261          Asia          Spratly

```

[262 rows x 2 columns]

```

[10]: newalcoholdata=pd.merge(countries, alcoholdata)
      newalcoholdata

```

```

[10]:      Continent      country  beer_servings  spirit_servings  \
0          Asia  Afghanistan             0             0
1        Europe    Albania             89            132
2        Africa    Algeria             25             0
3        Europe    Andorra            245            138
4        Africa    Angola            217             57
..          ...          ...          ...          ...
159         Asia    Uzbekistan             25            101
160  South America    Venezuela            333            100
161        Oceania      Samoa            105             18
162         Asia      Yemen              6             0
163        Africa    Zambia             32             19

```

```

      wine_servings  total_litres_of_pure_alcohol
0                0                0.0
1               54                4.9
2               14                0.7
3              312               12.4
4               45                5.9
..          ...          ...
159              8                2.4
160              3                7.7
161             24                2.6
162              0                0.1
163              4                2.5

```

[164 rows x 6 columns]

1.0.4 Data Cleaning: Removing duplicate countries from the dataframe

Given that we combined 2 different datasets, I want to make sure that there are no null values in the columns. I would do this using: - `df.isna().sum()`: Displays number of null values in each of the columns of the dataframe

What I can gather from this dataset is: - There are no null values so I will not need to remove any column / rows with null values. - There are 164 rows and 6 columns in this dataset - Within the continent column, there are 6 unique variable which are Asian, Europe, Africa, South America, Oceania, and North America. This poses a potential opportunity to use groupby to summarize our

findings. - Within the country column, there are 156 unique values which means that there are 8 countries in the column that have duplicated values. - `df.duplicated(subset=['col'])`: Displays boolean (T/F) values of duplicate rows. Keep determines which duplicates to keep (first = first occurrence is False default, last = last occurrence is False) - `df[df.duplicated(subset=['col'])]`: Displays dataframe of duplicate values - `df.drop_duplicates()`: If we want to keep the first occurrences, put `keep = first`, else = last, False= drop all duplicates

```
[11]: newalcoholdata.shape
```

```
[11]: (164, 6)
```

```
[12]: newalcoholdata.isna().sum()
```

```
[12]: Continent          0
      country            0
      beer_servings      0
      spirit_servings     0
      wine_servings      0
      total_litres_of_pure_alcohol  0
      dtype: int64
```

```
[13]: newalcoholdata['Continent'].nunique()
```

```
[13]: 6
```

```
[14]: newalcoholdata['Continent'].unique()
```

```
[14]: array(['Asia', 'Europe', 'Africa', 'South America', 'Oceania',
            'North America'], dtype=object)
```

```
[15]: newalcoholdata['country'].nunique()
```

```
[15]: 156
```

We see below that several Middle Eastern countries such as Azerbaijan, Armenia, Cyprus, Georgia, Kazakhstan, and Turkey show both Asia and Europe as the continent. This is because they span between Asian and Europe; they could be interpreted differently by people.

In order to keep the consistency, I will drop the second duplicate for all these values. I went with this decision based on the categorization of Netherlands since it is in Europe and not in North America. Now, we see that the new dataframe **newalcoholdata** has 156 rows and 6 columns, which is what we want. It has the same number of unique values in the 'country' column.

```
[16]: newalcoholdata[newalcoholdata.duplicated(subset=['country'], keep=False)]
```

```
[16]:
```

	Continent	country	beer_servings	spirit_servings	\
5	Europe	Azerbaijan	21	46	
6	Asia	Azerbaijan	21	46	

13	Europe	Armenia	21	179
14	Asia	Armenia	21	179
35	Africa	Congo	76	1
36	Africa	Congo	76	1
39	Europe	Cyprus	192	154
40	Asia	Cyprus	192	154
53	Europe	Georgia	52	100
54	Asia	Georgia	52	100
77	Europe	Kazakhstan	124	246
78	Asia	Kazakhstan	124	246
107	Europe	Netherlands	251	88
108	North America	Netherlands	251	88
149	Europe	Turkey	51	22
150	Asia	Turkey	51	22

	wine_servings	total_litres_of_pure_alcohol
5	5	1.3
6	5	1.3
13	11	3.8
14	11	3.8
35	9	1.7
36	9	1.7
39	113	8.2
40	113	8.2
53	149	5.4
54	149	5.4
77	12	6.8
78	12	6.8
107	190	9.4
108	190	9.4
149	7	1.4
150	7	1.4

```
[17]: newalcoholdata.drop_duplicates(subset=['country'], keep='first', inplace=True)
newalcoholdata.shape
```

```
[17]: (156, 6)
```

1.0.5 Data Grouping and Categorization: Retrieving summary statistics of the continents

Next, I wanted to analyze the difference in beer, spirit, and wine servings across the different continents and / or countries. I used the `groupby` function to segregate the data by continent and country. - `df.groupby(['col to group by'])[col to display].aggregate()`: This function allows us to split data into separate groups to better perform analysis

However, I noticed that the continent column, which had 6 unique variables (Asian, Europe, Africa, South America, Oceania, and North America) poses a potential opportunity to use pivot to sum-

marize our findings. - `pd.pivot_table(df, index='col', values='col', aggfunc='x')`: This function allows us to summarize data of a larger table

```
[18]: newalcoholdata.groupby(['Continent', 'country']).sum()
```

```
[18]:
```

		beer_servings	spirit_servings	wine_servings	\
Continent	country				
Africa	Algeria	25	0	14	
	Angola	217	57	45	
	Benin	34	4	13	
	Botswana	173	35	35	
	Burundi	88	0	0	
...		
South America	Paraguay	213	117	74	
	Peru	163	160	21	
	Suriname	128	178	7	
	Uruguay	115	35	220	
	Venezuela	333	100	3	
		total_litres_of_pure_alcohol			
Continent	country				
Africa	Algeria		0.7		
	Angola		5.9		
	Benin		1.1		
	Botswana		5.4		
	Burundi		6.3		
...			...		
South America	Paraguay		7.3		
	Peru		6.1		
	Suriname		5.6		
	Uruguay		6.6		
	Venezuela		7.7		

[156 rows x 4 columns]

Findings #1: - As shown in the **grouped bar chart**, out of the 6 continents, Europe has the highest number of alcohol servings across all types (beer, spirit, wine). Since there is one column that uses litres as a measurement unit, while the other 3 columns use number of servings, I will exclude the liters column for this visualization. - As shown in the **pie chart**, I wanted to ensure that this finding is not skewed by larger number of countries within Europe, so I decided to calculate the number of countries categorized within each continent using `df.groupby` to see if that affects the alcohol servings across continents. It turns out that Europe and Africa continents have the largest number of countries, followed by Asia, North America, South America, and Oceania. This shows that there is no strong correlation between alcohol servings and number of countries within each continent. Although the number of countries in Europe and Africa continents are the same, Europe has a significantly higher alcohol servings than Asia.

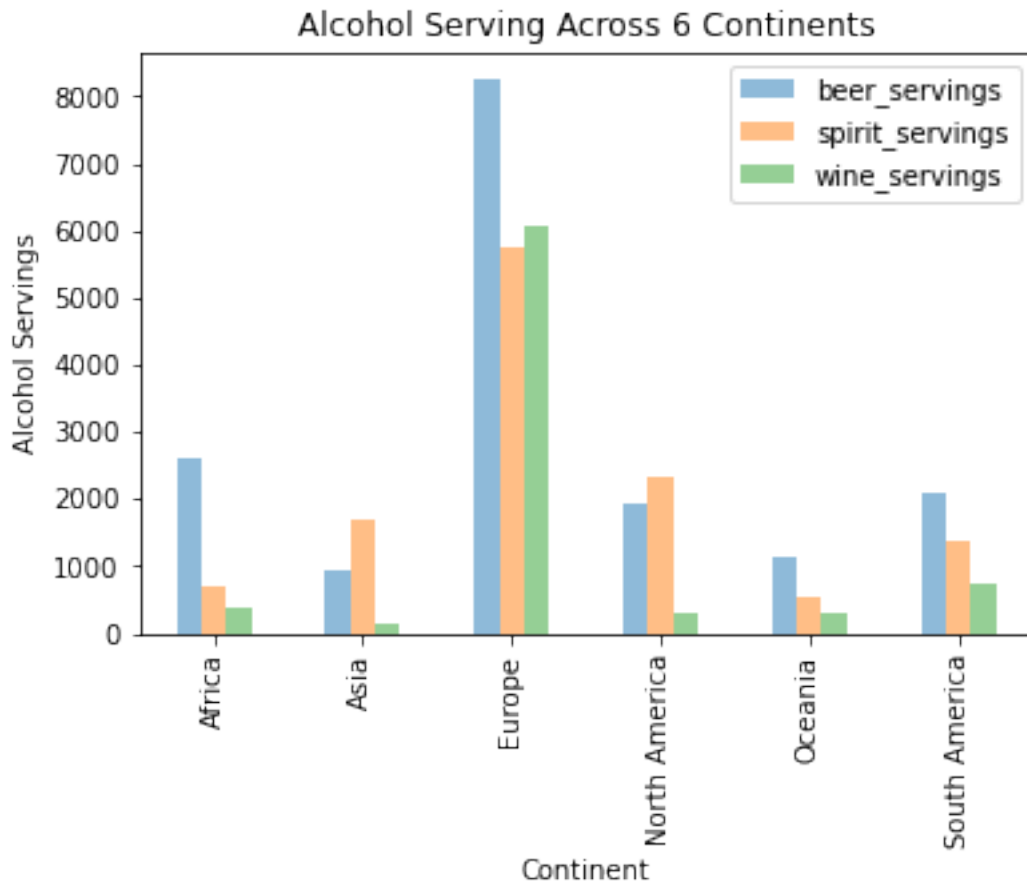

```
[19]: df_pivot=pd.pivot_table(newalcoholdata,
    ↪index='Continent',values=['beer_servings','spirit_servings','wine_servings'],
    ↪aggfunc='sum')
df_pivot
```

```
[19]:
```

	beer_servings	spirit_servings	wine_servings
Continent			
Africa	2605	680	374
Asia	938	1691	148
Europe	8239	5764	6082
North America	1918	2326	292
Oceania	1132	552	319
South America	2101	1377	749

```
[20]: #Creatng a grouped bar chart using matplotlib
ax=df_pivot.plot(kind="bar",alpha=0.5)
plt.title('Alcohol Serving Across 6 Continents')
plt.xlabel('Continent')
plt.ylabel('Alcohol Servings')

#Show plot
plt.show()
```




```
[21]: groupedcontinent=newalcoholdata.groupby('Continent')['country'].count()
groupedcontinent
```

```
[21]: Continent
Africa      43
Asia        33
Europe      43
North America  14
Oceania      11
South America  12
Name: country, dtype: int64
```

```
[22]: colors = ['lightsteelblue', 'thistle', 'pink', 'darkkhaki', 'wheat', 'lightcoral']
groupedcontinent.plot.pie(colors=colors)
```

```
[22]: <AxesSubplot:ylabel='country'>
```



Since we narrowed down that Europe has the largest alcohol consumption, I am interested in finding out which countries within Europe have the highest number of beer, spirit, and wine servings respectively.

- `df['col'].max()`: This shows us the max value
- `df['col']=`df['col'].max()`: This shows us true / false boolean values - `df[df['col']=`df['col'].max()]`: This returns the rows with the max column value
- `pd.concat([df1,df2,df3])`: This concatenates the dataframes vertically. Concat, merge and join function similarly. The difference is that *concat* combines dataframes *vertically*, whereas *merge and join* combines dataframes *horizontally*.

```
[58]: europedata=newalcoholdata[newalcoholdata['Continent']=='Europe']
      europedata.head()
```

```
[58]:   Continent    country  beer_servings  spirit_servings  wine_servings  \
1      Europe  Albania          89          132           54
3      Europe  Andorra         245          138          312
5      Europe  Azerbaijan         21           46           5
9      Europe   Austria         279           75          191
13     Europe   Armenia          21          179           11

      total_litres_of_pure_alcohol
1                                4.9
3                               12.4
5                                1.3
9                                9.7
13                               3.8
```

```
[54]: europedata[europedata['beer_servings']==europedata['beer_servings'].max()]
```

```
[54]:  Continent  country  beer_servings  spirit_servings  wine_servings  \
56      Europe  Germany             346             117             175

      total_litres_of_pure_alcohol
56                             11.3
```

```
[55]: europedata['beer_servings'].max()
```

```
[55]: 346
```

```
[61]: (europedata['beer_servings']==europedata['beer_servings'].max()).head()
```

```
[61]: 1      False
3      False
5      False
9      False
13     False
Name: beer_servings, dtype: bool
```

```
[74]: europebeer=europedata[europedata['beer_servings']==europedata['beer_servings'].
      ↪max()]
europebeer
```

```
[74]:  Continent  country  beer_servings  spirit_servings  wine_servings  \
56      Europe  Germany             346             117             175

      total_litres_of_pure_alcohol
56                             11.3
```

```
[77]: europewine=europedata[europedata['wine_servings']==europedata['wine_servings'].
      ↪max()]
europewine
```

```
[77]:  Continent  country  beer_servings  spirit_servings  wine_servings  \
50      Europe  France             127             151             370

      total_litres_of_pure_alcohol
50                             11.8
```

```
[76]: europespirit=europedata[europedata['spirit_servings']==europedata['spirit_servings'].
      ↪max()]
europespirit
```

```
[76]:  Continent  country  beer_servings  spirit_servings  wine_servings  \
26      Europe  Belarus             142             373             42
```

```

    total_litres_of_pure_alcohol
26                                14.4

```

```
[71]: europebeer.merge(europewine,on='Continent').merge(europespirit,on='Continent')
```

```
[71]:  Continent country_x  beer_servings_x  spirit_servings_x  wine_servings_x  \
0    Europe   Germany              346              117              175

    total_litres_of_pure_alcohol_x country_y  beer_servings_y  \
0                                11.3   France              127

    spirit_servings_y  wine_servings_y  total_litres_of_pure_alcohol_y  \
0                  151              370                  11.8

    country  beer_servings  spirit_servings  wine_servings  \
0  Belarus              142              373              42

    total_litres_of_pure_alcohol
0                                14.4

```

```
[79]: europemaxes=pd.concat([europebeer,europewine,europespirit])
europemaxes.reset_index()
```

```
[79]:   index  Continent  country  beer_servings  spirit_servings  wine_servings  \
0     56    Europe  Germany              346              117              175
1     50    Europe  France              127              151              370
2     26    Europe  Belarus              142              373              42

    total_litres_of_pure_alcohol
0                                11.3
1                                11.8
2                                14.4

```

On the flip side, I want to see which countries do not consume any alcohol. -
 (df['col1']=='x')&(df['col2']=='x'): This shows us a boolean value. If we put wrap this
 function with a `df[]` on the outside, it will return a dataframe.

If you want to filter one columnn for multiple criterias, this is a good [resource](#).

```
[96]: (newalcoholdata['wine_servings']==0)&(newalcoholdata['beer_servings']==0)
```

```
[96]: 0      True
1     False
2     False
3     False
4     False
...
```

```

159     False
160     False
161     False
162     False
163     False
Length: 156, dtype: bool

```

```

[97]: #filtering multiple columns by same criteria
noalcohol=newalcoholdata[((newalcoholdata['wine_servings']==0)&(newalcoholdata['beer_servings']
noalcohol

```

```

[97]:      Continent      country  beer_servings  spirit_servings  wine_servings  \
0         Asia  Afghanistan              0              0              0
12        Asia   Bangladesh              0              0              0
70        Asia         Iran              0              0              0
81        Asia      Kuwait              0              0              0
91        Asia   Maldives              0              0              0
94        Africa  Mauritania              0              0              0
97        Europe    Monaco              0              0              0
117       Asia    Pakistan              0              0              0
136       Africa    Somalia              0              0              0

      total_litres_of_pure_alcohol
0                                0.0
12                               0.0
70                               0.0
81                               0.0
91                               0.0
94                               0.0
97                               0.0
117                              0.0
136                              0.0

```

```

[98]: noalcohol.count()

```

```

[98]: Continent              9
country                    9
beer_servings              9
spirit_servings            9
wine_servings              9
total_litres_of_pure_alcohol  9
dtype: int64

```

1.0.6 Conclusion:

Europe has the largest consumption of alcohol across the world. Germany, France and Belarus takes the crown for beers, wine, and spirits respectively.

As for countries that do not consume alcohol, there are 9 countries in our list which of which most of them are in Asia. Based on this [article](#) I found, I believe it is because these countries have a sizable muslim population, therefore consumption of alcohol is prohibited.